

# High End Simulation in Practice: Molecular Dynamics Part 2

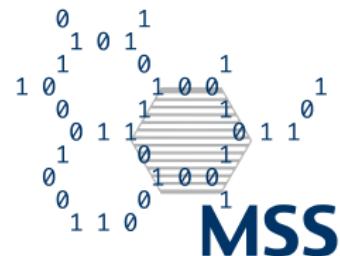
Severin Strobl

Institute for Multiscale Simulation

2014-05-07



FRIEDRICH-ALEXANDER  
UNIVERSITÄT  
ERLANGEN-NÜRNBERG  
TECHNISCHE FAKULTÄT



# Part 1: Neighbour Lists

## Recap

### Sorting Particles

Reducing Computational Costs

Short Range Forces

The Problem

The Idea

### Thread-Safe Implementation

The Solution

Reading the Data Structure

Conflicting Operations

Periodic Boundary Conditions

# Part 2: Smoothed Particle Hydrodynamics

## Introduction

## About Hydrodynamics

General Overview

Balance Equations

Solution Methods

Grids

Smoothing Function

## Resulting Force Laws

Selected Smoothing Function

Pressure Force

## Integrator

## Showroom

## Final Remarks

# Part I

## Neighbour Lists

# Outline

## Recap

### Sorting Particles

Reducing Computational Costs

Short Range Forces

The Problem

The Idea

### Thread-Safe Implementation

The Solution

Reading the Data Structure

Conflicting Operations

Periodic Boundary Conditions

- ▶ In the previous lecture, we learned how a simple MD program might work.
- ▶ Let us recap...

# Algorithm Outline

A simple molecular dynamics algorithm can be broken down as follows.

1. Take  $N$  particles, set initial positions (e.g. regular grid) and velocities (e.g. Gaussian distributed).
2. Calculate forces on the particles.

$$\mathbf{F}_i = \mathbf{F}_{body}(\mathbf{x}_i, \mathbf{v}_i) + \sum_{j \neq i} \mathbf{F}_2(\mathbf{x}_i, \mathbf{v}_i, \mathbf{x}_j, \mathbf{v}_j)$$

3. Integrate to the next time step.

$$\mathbf{x}_i(t + \Delta t) = \mathbf{x}_i(t) + \mathbf{v}_i(t)\Delta t + \frac{\mathbf{a}_i(t)\Delta t^2}{2}$$

$$\mathbf{v}_i(t + \Delta t) = \mathbf{v}_i(t) + \frac{\mathbf{a}_i(t) + \mathbf{a}_i(t + \Delta t)}{2}\Delta t$$

4. If we've simulated long enough, stop, otherwise goto 2.

# Force Models

- ▶ We had two examples for force laws in our last lecture:

$$\mathbf{F}_2^{LJ}(\mathbf{x}_i, \mathbf{x}_j) = 24 \varepsilon \left( \frac{\sigma}{x_{ij}} \right)^6 \left[ 2 \left( \frac{\sigma}{x_{ij}} \right)^6 - 1 \right] \frac{\mathbf{x}_{ij}}{x_{ij}^2}$$

$$\mathbf{F}_2^{SD}(\mathbf{x}_i, \mathbf{v}_i, \mathbf{x}_j, \mathbf{v}_j) = K \boldsymbol{\xi} + \gamma \dot{\boldsymbol{\xi}}$$

- ▶ Another force commonly used for granular matter is the viscoelastic Hertz model:

$$\mathbf{F}_2^{Hertz}(\mathbf{x}_i, \mathbf{v}_i, \mathbf{x}_j, \mathbf{v}_j) = \frac{2Y\sqrt{R_{eff}}}{3(1-\nu^2)} \left( \sqrt{\boldsymbol{\xi} \boldsymbol{\xi}} + \frac{3}{2} A \sqrt{\boldsymbol{\xi} \dot{\boldsymbol{\xi}}} \right),$$

with  $Y$ ,  $\nu$  and  $R_{eff}$  denoting the Young's modulus, the Poisson's ratio and the effective radius, respectively.  $A$  is a dissipative constant, controlling how much energy is transferred into internal energy in each collision.

# Outline

## Recap

## Sorting Particles

Reducing Computational Costs

Short Range Forces

The Problem

The Idea

## Thread-Safe Implementation

The Solution

Reading the Data Structure

Conflicting Operations

Periodic Boundary Conditions

## Algorithmic Complexity

Let's have a look at the computational complexity of the different steps in a MD simulation:

1.  $\mathcal{O}(N)$ : Take  $N$  particles, on a regular grid with Gaussian distributed velocities.
2.  $\mathcal{O}(N^2)$ : Calculate forces on the particles.

$$\mathbf{F}_i = \mathbf{F}_{body}(\mathbf{x}_i, \mathbf{v}_i) + \sum_{j \neq i} \mathbf{F}_2(\mathbf{x}_i, \mathbf{v}_i, \mathbf{x}_j, \mathbf{v}_j)$$

3.  $\mathcal{O}(N)$ : Integrate to the next time step.

$$\mathbf{x}_i(t + \Delta t) = \mathbf{x}_i(t) + \mathbf{v}_i(t)\Delta t + \frac{\mathbf{a}_i(t)\Delta t^2}{2}$$

$$\mathbf{v}_i(t + \Delta t) = \mathbf{v}_i(t) + \frac{\mathbf{a}_i(t) + \mathbf{a}_i(t + \Delta t)}{2}\Delta t$$

4. If we've simulated long enough, stop, otherwise goto 2.

# Algorithmic Complexity

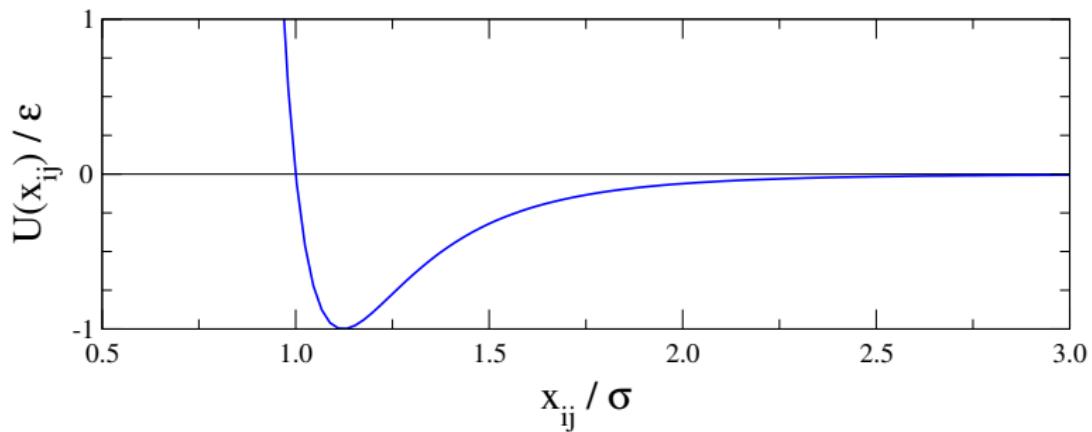
2.  $\mathcal{O}(N^2)$ : Calculate forces on the particles.

$$\mathbf{F}_i = \mathbf{F}_{body}(\mathbf{x}_i, \mathbf{v}_i) + \sum_{j \neq i} \mathbf{F}_2(\mathbf{x}_i, \mathbf{v}_i, \mathbf{x}_j, \mathbf{v}_j)$$

- ▶ The force calculation is by far the most expensive step.
- ▶ It scales as  $\mathcal{O}(N^2)$ , as we sum  $N$  terms for each of the  $N$  particles.
- ▶ How can we reduce this cost?

## Short Range Forces

- ▶ Lets take a look at the Lennard-Jones potential again.



- ▶ The gradient of the potential is pretty shallow for larger distances  $x_{ij}$ .
- ▶ Therefore the force tends towards zero with increasing separation distance.

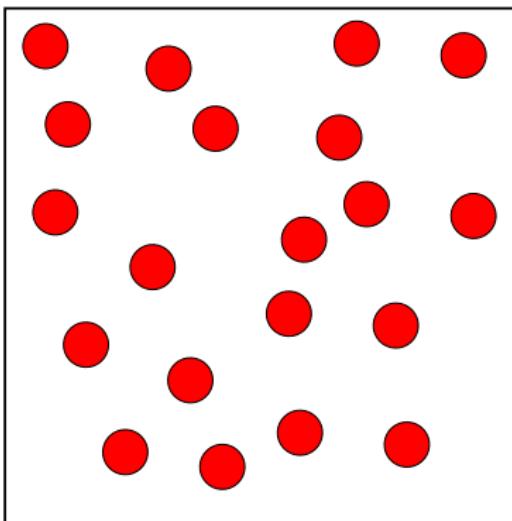
## Short Range Forces

- ▶ We can then approximate the Lennard-Jones force term by truncating at a certain distance  $r_c$ :

$$\mathbf{F}_{2, \text{trunc}}^{LJ}(\mathbf{x}_i, \mathbf{x}_j) = \begin{cases} 24 \varepsilon \left( \frac{\sigma}{x_{ij}} \right)^6 \left[ 2 \left( \frac{\sigma}{x_{ij}} \right)^6 - 1 \right] \frac{\mathbf{x}_{ij}}{x_{ij}^2} & \text{for } x_{ij} \leq r_c \\ 0 & \text{for } x_{ij} > r_c \end{cases}$$

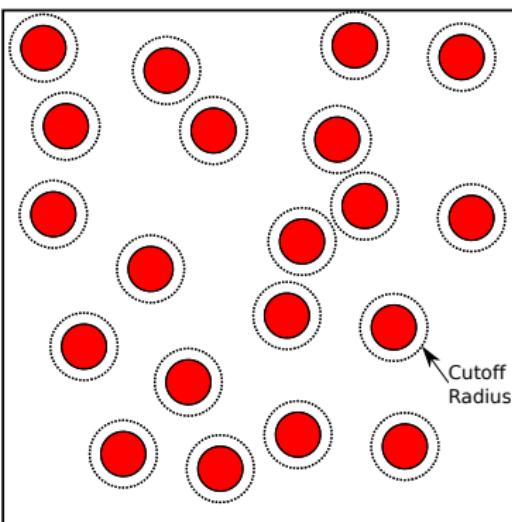
- ▶ This distance  $r_c$  where the potential is truncated is commonly referred to as the *cut-off distance* or *cut-off radius*.
- ▶ A popular choice for  $r_c$  is  $2.5\sigma$ , which is a good compromise between computational cost and the error introduced by the truncation of the potential.
- ▶ When truncating the potential, we only need to sum up the forces between neighboring particles.
- ▶ So now we need a way to find neighboring particles quickly, we need a *neighbor list*!

## The System



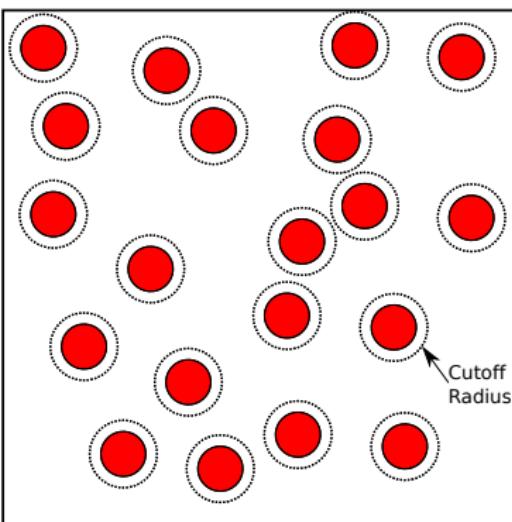
- ▶ We have a box of particles.

# The System



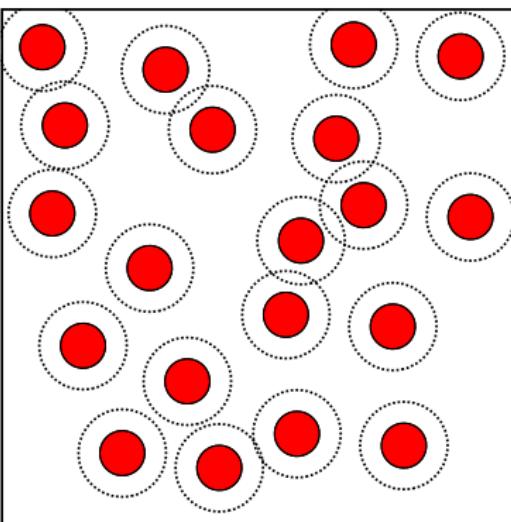
- ▶ These particles only have short range interactions with each other.
- ▶ Before we performed the force summations by checking all particles against each other.

## The System



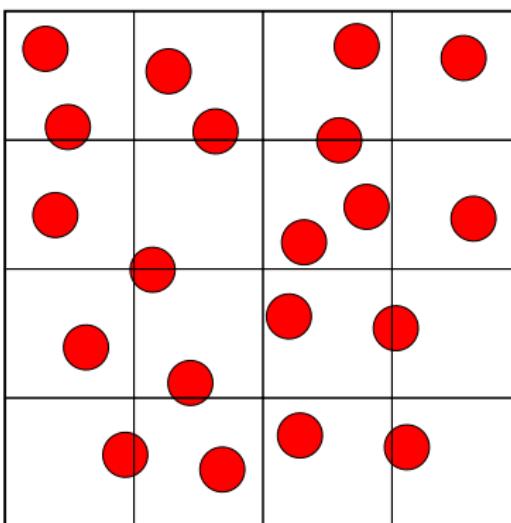
- ▶ But this is a huge waste of computation for short range forces.
- ▶ If the interactions were as short ranged as in the image, no forces would have to be computed at all!

## The System



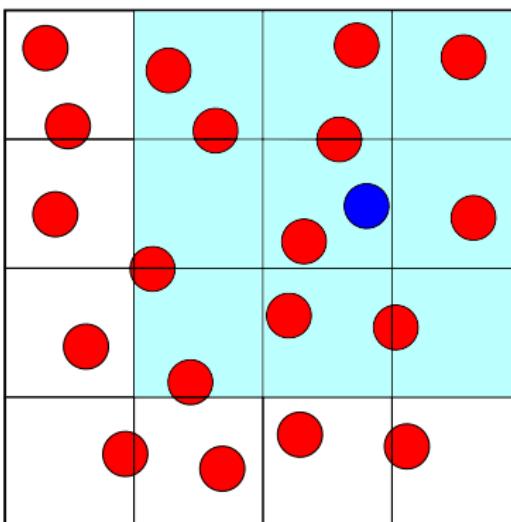
- ▶ Typically they're not quite so short ranged.
- ▶ But using a neighbor list reduces the force calculations from  $O(N^2)$  to  $O(N) \times N_{neighbors}$ , so it almost always pays off.

## The Idea



- ▶ We start by sorting the particles into a grid.
  - ▶ This lets us find neighbouring (localised) particles quickly.

## The Idea



- ▶ The width of a single cell must be  $\geq r_{cutoff}$ .
- ▶ Therefore, all particles within  $r_{cutoff}$  of a single particle (blue) are also within a  $3 \times 3 \times 3$  group of cells.

# Outline

## Recap

### Sorting Particles

Reducing Computational Costs

Short Range Forces

The Problem

The Idea

### Thread-Safe Implementation

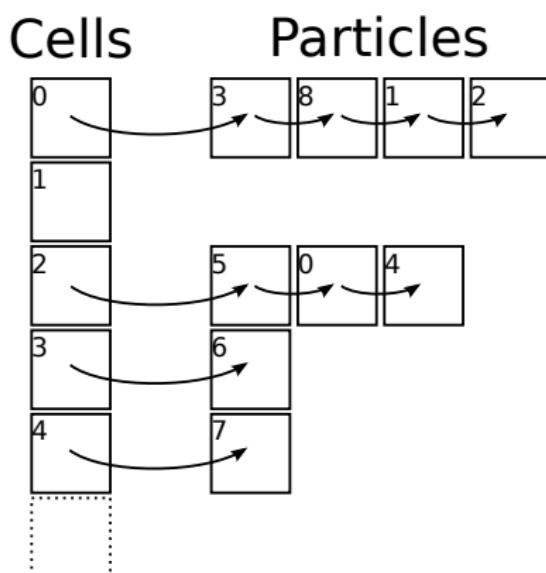
The Solution

Reading the Data Structure

Conflicting Operations

Periodic Boundary Conditions

# The Idea



- ▶ The particles in each cell are stored in a linked list.
- ▶ These lists are updated as particles move from cell to cell.

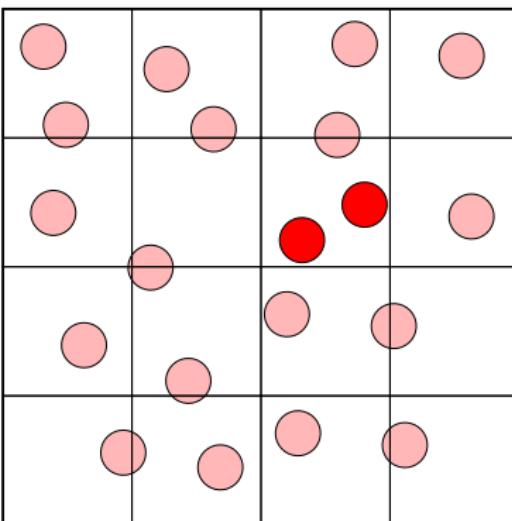
# The Idea

- ▶ Finding the cell a particle belongs to is easy: division and rounding to integer.
- ▶ Adding a particle to a cell is difficult in a parallel environment like OpenCL:
  - ▶ The list of particles belonging to a cell must be available to all threads at the same time and can be large.
  - ▶ But what if multiple threads try to add a particle to the same cell at the same time? Particles might be lost, data structures corrupted,...
- ▶ Atomic operations are here to save the day!

# The Solution

- ▶ How do we perform this sorting into cells using multiple threads?
- ▶ Do we make cells look for their contained particles? ( $\mathcal{O}(N^2)$ )
- ▶ Do we make particles look for their containing cells? ( $\mathcal{O}(N)$ , but requires locks or atomic operations)
- ▶ Here we'll do the latter as it is  $\mathcal{O}(N)$ , even with the slowdown of atomic operations it will outperform the first option even at moderate system sizes.

# The Solution



By only considering the particles center point we can state:

- ▶ Each particle belongs to only one cell.
- ▶ Cells can contain multiple particles.

# The Data Structure

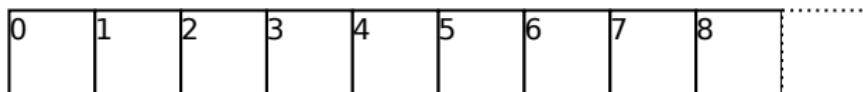
- ▶ These statements allow us to choose our data structure.
- ▶ We could use a singly linked list to store the contents of a single cell.
- ▶ Cells are assigned IDs from 0 to  $N_{cells} - 1$ .
- ▶ Particles are assigned IDs from 0 to  $N - 1$ .
- ▶ An integer is then needed for each particle and cell.

# The Data Structure

Cells



Particles



- ▶ Two arrays in memory hold the linked list integers of the cell and the particles.

# The Data Structure

Cells

0	-1	-1	-1	-1	-1	-1	-1	-1	-1	...
---	----	----	----	----	----	----	----	----	----	-----

Particles

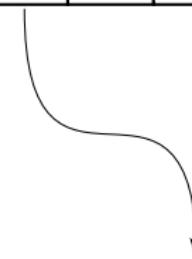
0	1	2	3	4	5	6	7	8	...
---	---	---	---	---	---	---	---	---	-----

- ▶ The cells are initialized to an empty state (-1).

# The Data Structure

Cells

0	-1	-1	2		3	-1	-1	5	-1	6	-1	7	-1	-1	-1		
---	----	----	---	--	---	----	----	---	----	---	----	---	----	----	----	--	--



Particles

0	1	2	3	4	-1	5	6	7	8							
---	---	---	---	---	----	---	---	---	---	--	--	--	--	--	--	--

- ▶ We want to add particle 4 to cell number 2.
- ▶ We move the value of cell 2 to particle 2's integer.

# The Data Structure

Cells

0	-1	-1	2	4	3	-1	-1	5	-1	6	-1	7	-1	8	-1	...	...
---	----	----	---	---	---	----	----	---	----	---	----	---	----	---	----	-----	-----



Particles

0	1	2	3	4	-1	5	6	7	8	...	...
---	---	---	---	---	----	---	---	---	---	-----	-----

- ▶ Now we write the ID of particle 4 in the integer of cell 2.
- ▶ Particle 4 is now added to cell 2.

# The Data Structure

Cells

0	1	2	0	3	4	5	6	7	8	...
-1	-1	0	-1	-1	-1	-1	-1	-1	-1	...

Particles

0	1	2	3	4	5	6	7	8	...
4				-1					...

- ▶ This can continue and multiple particles can be added to a single cell.
- ▶ Essentially, each particle's number is pushed onto the start of a cell's singly linked list.

# The Data Structure

Cells

0	1	2	0	3	4	5	6	7	8	...
-1	-1	0	-1	-1	-1	-1	-1	-1	-1	...



Particles

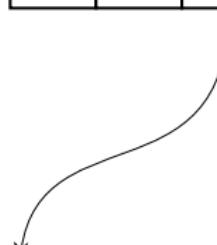
0	1	2	3	4	5	6	7	8	...
4				-1					...

- ▶ Adding a particle is a simple exchange procedure, but what about reading a cell's contents?

## The Data Structure

Cells

0	-1	1	-1	2	0	3	-1	4	-1	5	-1	6	-1	7	-1	8	-1	...	...
---	----	---	----	---	---	---	----	---	----	---	----	---	----	---	----	---	----	-----	-----



Particles

0	4	1		2		3		4	-1	5		6		7		8		...	...
---	---	---	--	---	--	---	--	---	----	---	--	---	--	---	--	---	--	-----	-----

- ▶ To read what particles are in cell 2, we check its integer for the first particle (0).

# The Data Structure

Cells

0	1	2	0	3	4	5	6	7	8	...
-1	-1	0	-1	-1	-1	-1	-1	-1	-1	...

Particles

0	1	2	3	4	5	6	7	8	...
4				-1					...

- ▶ The integer of the first particle (0), then leads us to the integer of particle 4.
- ▶ As soon as we encounter a -1, the end of the list has been reached.

# The Data Structure

Cells

0	-1	-1	2	3	-1	4	5	6	7	8	-1	...
---	----	----	---	---	----	---	---	---	---	---	----	-----

Particles

0	4		2	3	4	-1	5	4	6	7	8	...
---	---	--	---	---	---	----	---	---	---	---	---	-----

- ▶ A problem arises if multiple threads try to add a particle to the same cell.
- ▶ Their exchange operations might overlap, leading to an invalid state.

# The Data Structure

Cells

0	1	2	3	4	5	6	7	8	...
-1	-1	5	-1	-1	-1	-1	-1	-1	...

Particles

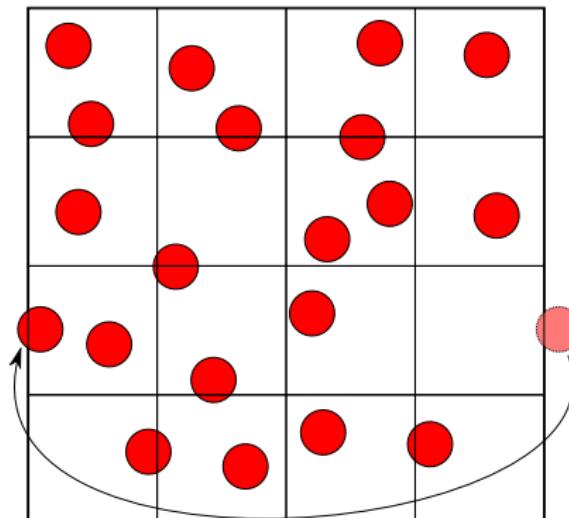
0	1	2	3	4	5	6	7	8	...
4				-1	0				...

atomic\_xchg

atomic\_xchg

- ▶ We solve this by making the exchange operation “atomic” (`atomic_xchg` in OpenCL).
- ▶ This ensures that our linked list is always valid.

## Final Notes



- ▶ For our kernel to work, all our particles must be within the gridded area, so we have to ensure they stay within the system boundaries.
- ▶ We use periodic boundary conditions for this and every time step we must ensure the particles are within the primary image.

# Final Notes

- ▶ This neighbor list approach allows our simulation to scale to much larger system sizes.
- ▶ There are still some speed-ups to be gained:
  - ▶ You can make the neighborhood area larger than needed and then the neighbor lists can remain valid for multiple time steps.
  - ▶ The cell neighbor list can be used to create a spherical neighbor list to reduce the number of neighbors even further.

## Part II

# Smoothed Particle Hydrodynamics

# Outline

## Introduction

### About Hydrodynamics

General Overview

Balance Equations

Solution Methods

Grids

Smoothing Function

### Resulting Force Laws

Selected Smoothing Function

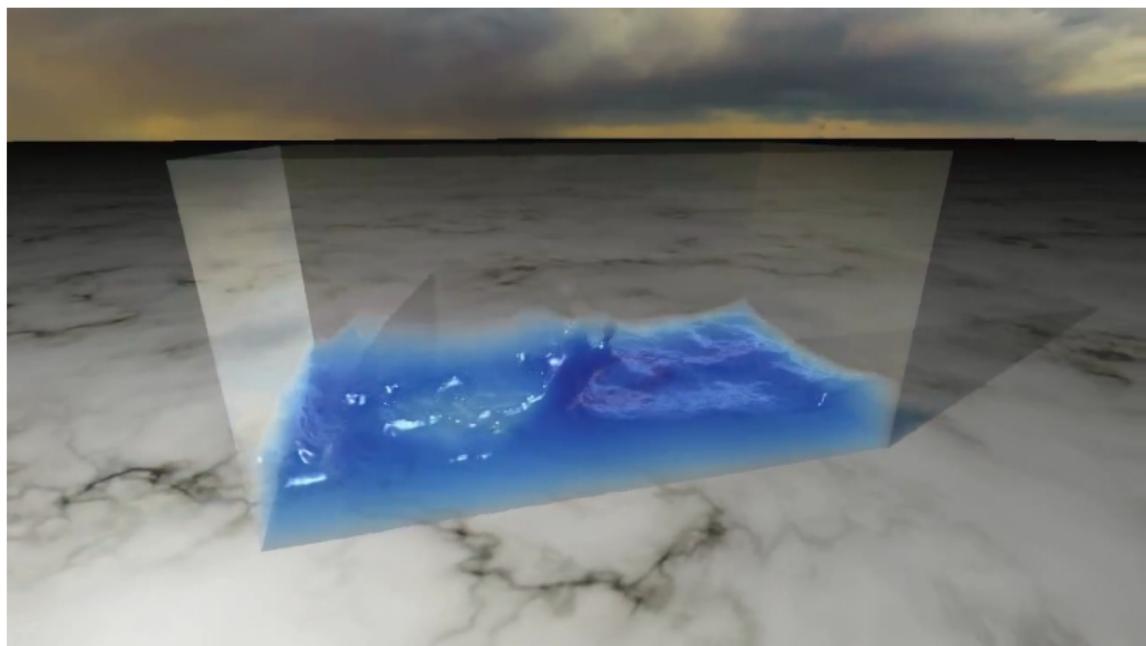
Pressure Force

## Integrator

## Showroom

## Final Remarks

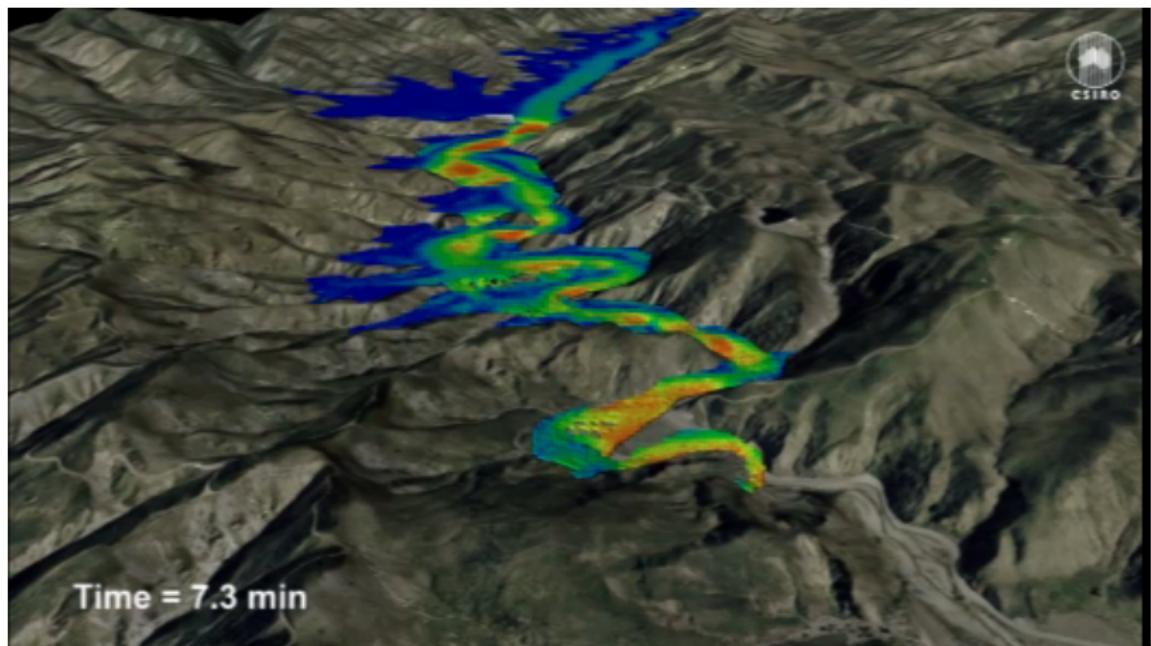
# Applications for Smoothed Particle Hydrodynamics (SPH)



Source:

NVIDIA PhysX Fluids Demo running on a GF100

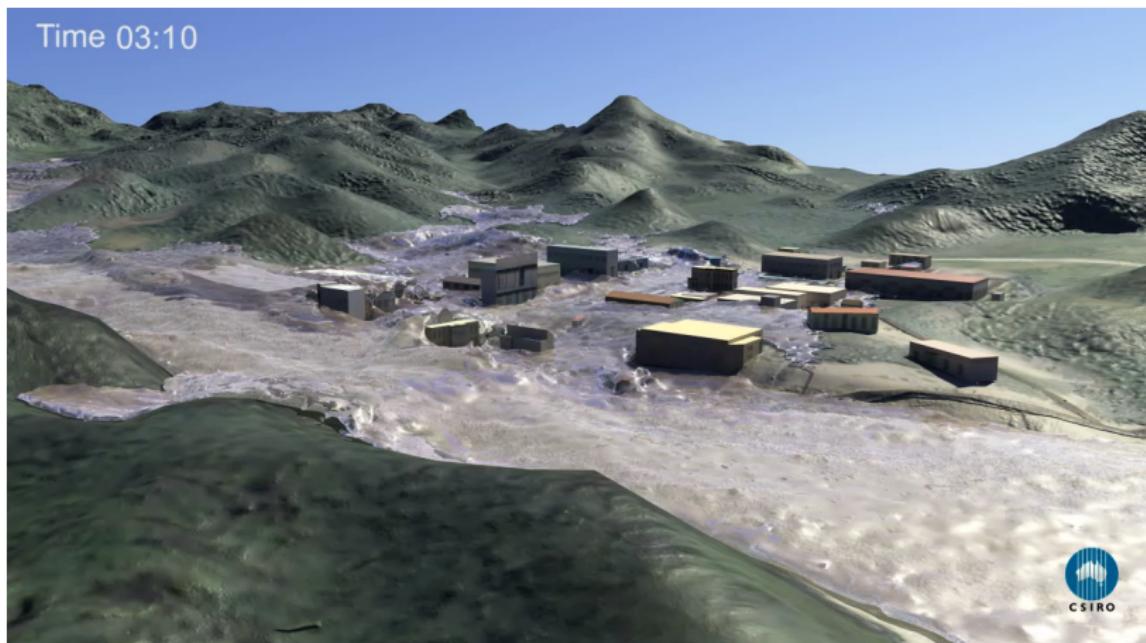
# Applications for Smoothed Particle Hydrodynamics (SPH)



Source:

CSIRO Mathematics, Informatics &amp; Statistics, Clayton, AU

# Applications for Smoothed Particle Hydrodynamics (SPH)



# Smoothed Particle Hydrodynamics (SPH)

- ▶ We're going to describe SPH in a rough and general fashion, as the underlying theory is not in the scope of this lecture.
- ▶ You can read several more thorough derivations on the internet or in one of the many books on the subject.
- ▶ This description is not complete and misses much/all of the necessary approximations required to derive SPH, so **Buyer Beware**.
- ▶ The presentation of the SPH method is based on the work of T. Harada et al.: Smoothed Particle Hydrodynamics on GPUs, 2007.

# Outline

## Introduction

### About Hydrodynamics

General Overview

Balance Equations

Solution Methods

Grids

Smoothing Function

## Resulting Force Laws

Selected Smoothing Function

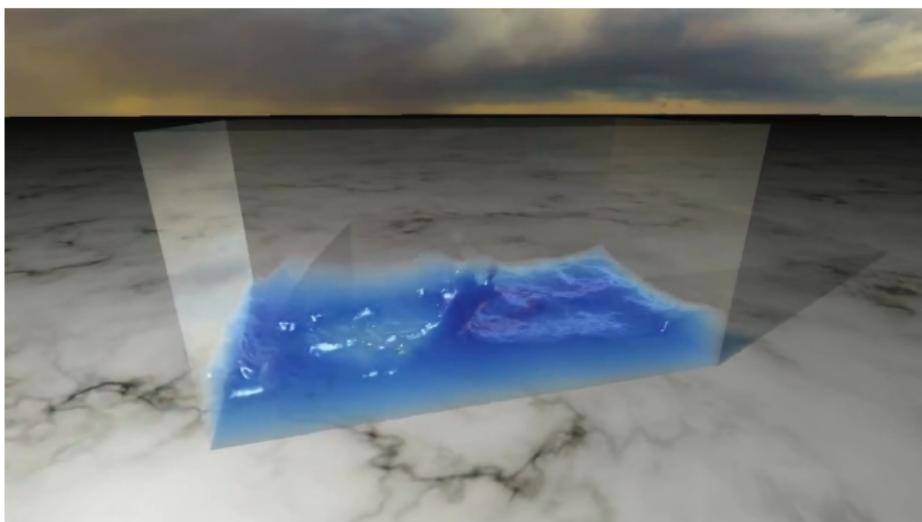
Pressure Force

## Integrator

## Showroom

## Final Remarks

# Smoothed Particle Hydrodynamics



- ▶ Smoothed Particle Hydrodynamics is a numerical method to solve a set of **hydrodynamic** equations.
- ▶ **Hydrodynamics** is the theory describing the motion of fluids.

# Smoothed Particle Hydrodynamics

- ▶ Hydrodynamics generally solves balance equations for the density, momentum and energy.
  - ▶ For example, the conservation of mass gives the continuity equation:

$$\frac{\partial \rho}{\partial t} + \nabla \cdot (\rho \mathbf{v}) = 0$$

- ▶ This equation simply states, that for an **element** (small volume) of fluid to change its **density**  $\rho$  (lose or gain some fluid and by that change its mass), that mass must flow **out** or flow **in** from/to the element.
- ▶ But we have two unknowns, and only one equation. In addition we need an expression for the **flow velocity**  $\mathbf{v}$ .

# Smoothed Particle Hydrodynamics

- ▶ Hydrodynamics generally concerns solving balance equations for the density, momentum and energy.
  - ▶ A momentum balance for the fluid element helps us here:

$$\rho \frac{D\mathbf{v}}{Dt} = \rho \left( \frac{\partial \mathbf{v}}{\partial t} + \mathbf{v} \cdot \nabla \mathbf{v} \right) = -\nabla p + \nabla \cdot \mathbf{T} + \mathbf{F}_{\text{other}},$$

with  $p$  denoting the pressure,  $\mathbf{T}$  being the deviatoric stress tensor and  $\mathbf{F}_{\text{other}}$  representing body forces acting on the fluid element (like gravity).

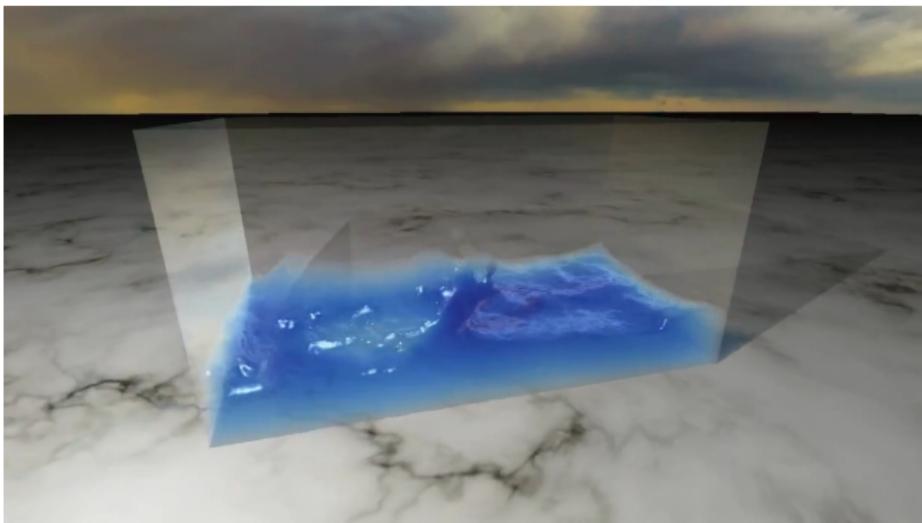
- ▶ This bears more than a passing resemblance to:

$$m \mathbf{a} = \mathbf{F} = \mathbf{F}_{\text{pressure}} + \mathbf{F}_{\text{viscous}} + \mathbf{F}_{\text{other}}$$

As they should, fluids obey Newton's equation of motion just like everything else (e.g. atoms in molecular dynamics).

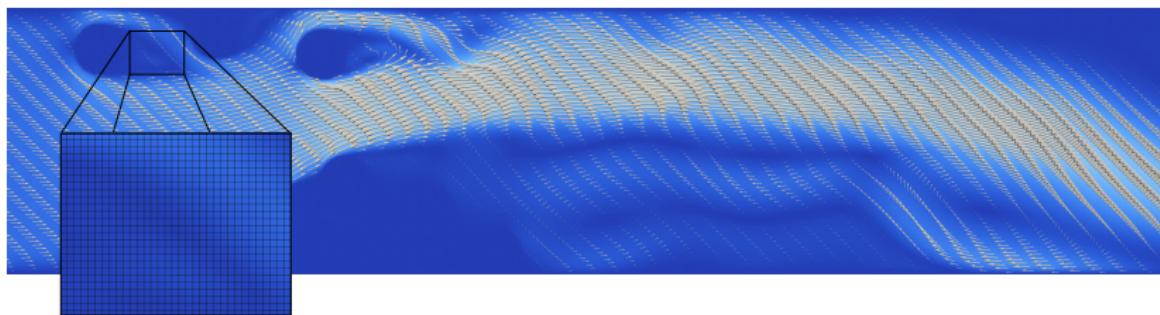
- ▶ For an incompressible Newtonian fluid the stress term  $\nabla \cdot \mathbf{T}$  simplifies to  $\mu \nabla^2 \mathbf{v}$ , so we get  $\mathbf{F}_{\text{viscous}} = \mu \nabla^2 \mathbf{v}$ , with  $\mu$  being the viscosity.

# Smoothed Particle Hydrodynamics



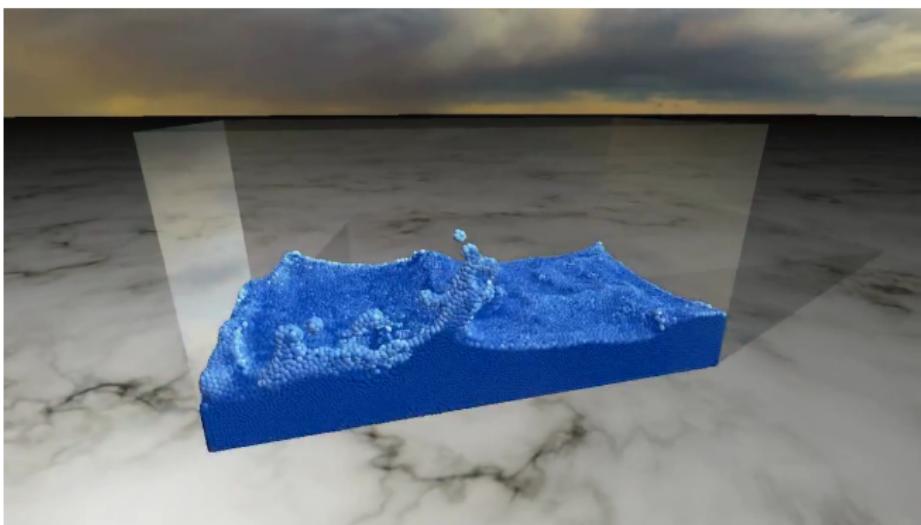
- ▶ Mathematically speaking these equations are only valid for a point of fluid (a differential element).
- ▶ But we have a finite volume of fluid to simulate and numerically we can only handle a finite number of “fluid points” or elements.

# Smoothed Particle Hydrodynamics



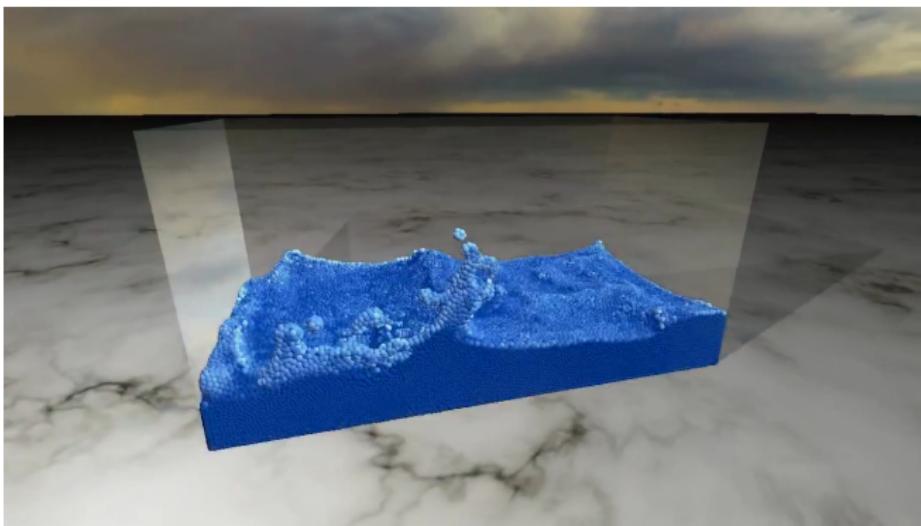
- ▶ A common approach is to use a grid of nodes to approximate the fluid.
- ▶ The shape and topology of this grid is a huge area of research as it must be:
  - ▶ Fine enough (many nodes) to capture all the essential behavior.
  - ▶ Coarse enough (few nodes) to have a manageable number of nodes and remain efficient.
- ▶ Only on the nodes of the grid are the hydrodynamic variables calculated or stored (density, temperature, pressure, velocity).

# Smoothed Particle Hydrodynamics



- ▶ SPH still utilizes a grid, but the nodes of the grid **move with the fluid**.
- ▶ Each node of the grid represents a finite mass of fluid.

# Smoothed Particle Hydrodynamics

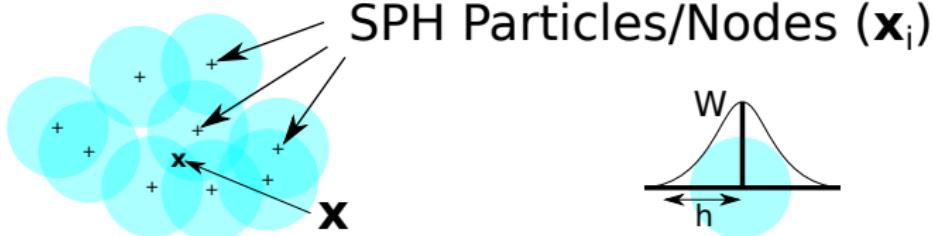


- ▶ This is the “Particle” part in the name of Smoothed Particle Hydrodynamics.
- ▶ We have particles of fluid, each with a mass and interacting according to the hydrodynamic equations.

# Smoothed Particle Hydrodynamics

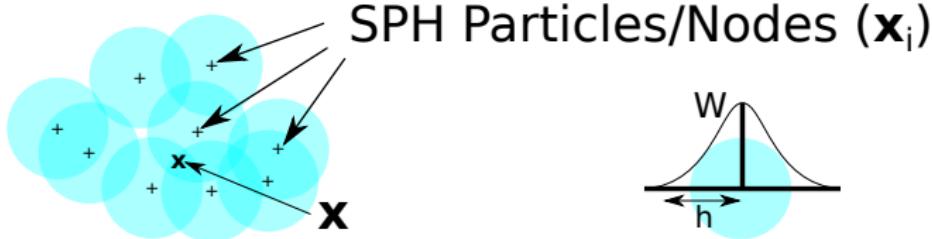
- ▶ A problem with this SPH grid is that the neighbors of each node/particle of the grid change continuously.
- ▶ With a fixed grid, we could determine spatial derivatives (needed for the fluxes) through looking at the well defined neighboring nodes.
- ▶ But what are the neighboring nodes of a constantly shifting grid?
- ▶ And nodes must be correctly weighted too for the spatial derivative, depending on orientation and distance!

## Smoothed Particle Hydrodynamics



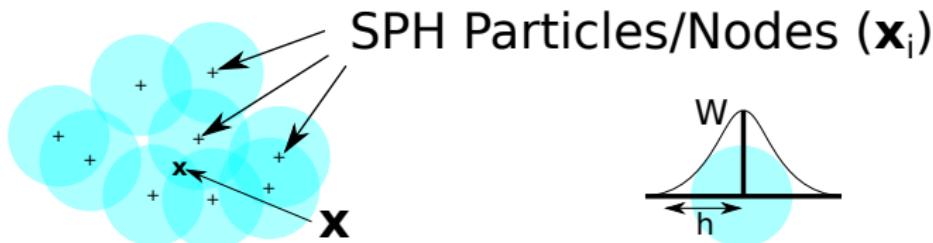
- ▶ To combat this difficulty, we will use a smoothing function  $W(\mathbf{x} - \mathbf{x}_i, h)$  to map our “particle” description to continuous hydrodynamic fields.
- ▶ This weights the contribution of nearby particles,  $\mathbf{x}_i$ , to the hydrodynamic properties at some point  $\mathbf{x}$ .

## Smoothed Particle Hydrodynamics



- ▶ This smoothing function  $W(\mathbf{x} - \mathbf{x}_i, h)$  also has a smoothing length,  $h$ , which describes how near particles have to be to contribute.
- ▶ This is the “Smoothed” part in the name of Smoothed Particle Hydrodynamics.

# Smoothed Particle Hydrodynamics



- ▶ The smoothing function is both a strength and a weakness of SPH.
- ▶ We don't have a grid to maintain, the fluid naturally clusters grid points near dense regions. We can even have free surface boundaries!
- ▶ But this smoothing function is also a curse for low density fluids, our approximation breaks down as SPH particles don't overlap. Our low density regions will have either 0 or small spikes in the density!

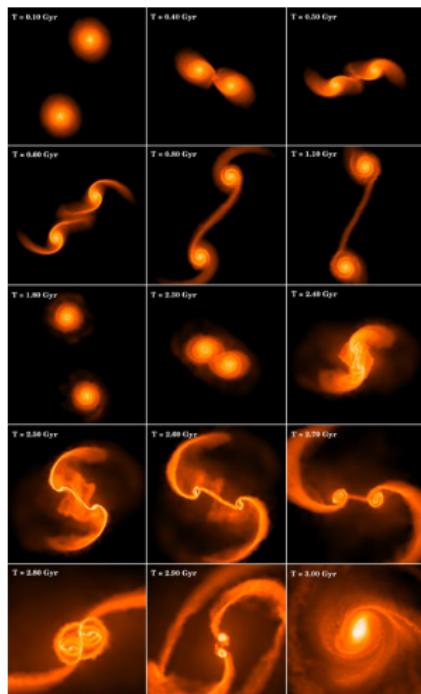


Image by Stelios Kazantzidis

[http://kicp.uchicago.edu/research/highlights/highlight\\_2006-03-21.html](http://kicp.uchicago.edu/research/highlights/highlight_2006-03-21.html)

- ▶ This can be counteracted by varying the smoothing length depending on the local density, but we won't cover this here.
- ▶ Varying the smoothing length is crucial in galactic simulations, a popular application of SPH.

## Smoothed Particle Hydrodynamics

- With a smoothing function we can define “smoothed” hydrodynamic properties

$$\langle A(\mathbf{x}) \rangle = \int A(\tau) W(\mathbf{x} - \tau, h) d\tau$$

where  $A$  is a hydrodynamic property, like the density.

- As our smoothing length vanishes ( $\lim_{h \rightarrow 0}$ ) and our number of SPH particles increases, our “smoothed” properties approach the true hydrodynamic properties.

# Smoothed Particle Hydrodynamics

- ▶ As the fluid is represented by particles, the previous integral over all space can be reduced to a sum over all particles:

$$\langle A(\mathbf{x}) \rangle = \sum_i \frac{m_i}{\rho_i} A(\mathbf{x}_i) W(\mathbf{x} - \mathbf{x}_i, h),$$

where  $A(\mathbf{x}_i)$  is the property of particle  $i$  (like the mass).

- ▶ This equation lets us calculate any property, at any point, using only the data at each particle's position.
- ▶ For example, as the mass of a single SPH particle is defined, the smoothed density of the fluid at any given point is given by:

$$\langle \rho(\mathbf{x}) \rangle = \sum_i \frac{m_i}{\rho_i} \rho_i W(\mathbf{x} - \mathbf{x}_i, h) = \sum_i m_i W(\mathbf{x} - \mathbf{x}_i, h)$$

# Smoothed Particle Hydrodynamics

- ▶ Even more crucially, spatial derivatives in a property (like the pressure term  $\mathbf{F}_{pressure} = -\nabla p$ ), can be rewritten in terms of gradients of the smoothing function!

$$\langle \nabla A(\mathbf{x}) \rangle = \sum_i \frac{m_i}{\rho_i} A(\mathbf{x}_i) \nabla W(\mathbf{x} - \mathbf{x}_i, h)$$

- ▶ Thus we can solve our hydrodynamic equations using SPH, we just need an appropriate smoothing function.
- ▶ If we solve the momentum balance equation at the point of an SPH particle, then we know how the particle moves (the particle follows the flow)!

# Outline

Introduction

About Hydrodynamics

General Overview

Balance Equations

Solution Methods

Grids

Smoothing Function

Resulting Force Laws

Selected Smoothing Function

Pressure Force

Integrator

Showroom

Final Remarks

# Smoothed Particle Hydrodynamics

- We want a spherical smoothing function that has the properties

$$\int W(\mathbf{x}, h) d\mathbf{x} = 1 \quad \text{and}$$

$$\lim_{h \rightarrow 0} W(\mathbf{x} - \mathbf{x}_i, h) = \delta(\mathbf{x} - \mathbf{x}_i)$$

- We could use a Gaussian, but this would mean that all particles would contribute to all points.
- Instead, we want a localized function that is zero if  $|\mathbf{x} - \mathbf{x}_i| \geq r_{cutoff} = h$  (for the neighbor lists).
- Often people use three piece splines.
- Here, we'll use the approach from the paper by T. Harada et al. and use the following function:

$$W(\mathbf{x} - \mathbf{x}_i, h) = \begin{cases} \frac{315}{64\pi h^9} \left( h^2 - |\mathbf{x} - \mathbf{x}_i|^2 \right)^3 & \text{for } |\mathbf{x} - \mathbf{x}_i| \leq h \\ 0 & \text{for } |\mathbf{x} - \mathbf{x}_i| > h \end{cases}$$

# Smoothed Particle Hydrodynamics

- ▶ For terms including gradients (e.g. pressure force, a function of the gradient of the pressure) we need the gradient of the smoothing function.
- ▶ For the pressure term we can use the gradient of a modified smoothing kernel (for details see M. Müller et al.: Particle-based Fluid Simulation for Interactive Applications, 2003):

$$\nabla W_p(\mathbf{x} - \mathbf{x}_i, h) = \begin{cases} -\frac{45}{\pi h^6} (h - |\mathbf{x} - \mathbf{x}_i|)^2 \frac{\mathbf{x} - \mathbf{x}_i}{|\mathbf{x} - \mathbf{x}_i|} & \text{for } |\mathbf{x} - \mathbf{x}_i| \leq h \\ 0 & \text{for } |\mathbf{x} - \mathbf{x}_i| > h \end{cases}$$

# Smoothed Particle Hydrodynamics

- ▶ The pressure force is then given by

$$\mathbf{F}_i^{\text{pressure}} = -\langle \nabla p(\mathbf{x}_i) \rangle = -\sum_j m_j \frac{p_i + p_j}{2\rho_j} \nabla W_p(\mathbf{x}_i - \mathbf{x}_j, h)$$

- ▶ To ensure that momentum is always conserved, this expression was made symmetric in  $i$  and  $j$ .
- ▶ In order to increase the numerical stability of the SPH method, the pressure at the location of a particle is not calculated via the ideal gas equation of state as

$$p = \rho k_B T = \rho k,$$

with  $k_B$  being the gas constant and  $T$  the temperature, but instead as

$$p = (\rho - \rho_0)k_B T = (\rho - \rho_0)k$$

The offset  $\rho_0$  is the rest density and mathematically has no effect on the pressure forces.

# Smoothed Particle Hydrodynamics

- ▶ We would like a viscous term too, and we define this now.
- ▶ The expression for the viscous force with the necessary modifications to make it symmetric is:

$$\mathbf{F}_i^{viscous} = \mu \sum_j m_j \frac{\mathbf{v}_j - \mathbf{v}_i}{\rho_j} \nabla^2 W_v(\mathbf{x}_i - \mathbf{x}_j, h)$$

- ▶ Where we use  $\nabla^2 W_v(\mathbf{x} - \mathbf{x}_i, h)$  according to M. Müller et al.:

$$\nabla^2 W_v(\mathbf{x} - \mathbf{x}_i, h) = \begin{cases} \frac{45}{\pi h^6} (h - |\mathbf{x} - \mathbf{x}_i|) & \text{for } |\mathbf{x} - \mathbf{x}_i| \leq h \\ 0 & \text{for } |\mathbf{x} - \mathbf{x}_i| > h \end{cases}$$

# Smoothed Particle Hydrodynamics

- ▶ This completes our inter-particle, hydrodynamic force definitions.
- ▶ Of course other (external) forces like gravity can also be added.
- ▶ We then need to integrate our system forward in time, just like the MD example...

# Outline

Introduction

About Hydrodynamics

General Overview

Balance Equations

Solution Methods

Grids

Smoothing Function

Resulting Force Laws

Selected Smoothing Function

Pressure Force

Integrator

Showroom

Final Remarks

# Smoothed Particle Hydrodynamics

- ▶ We'll only use a simple Euler integrator to evolve the system forward in time as the system is rather dissipative anyway

$$\mathbf{r}_i(t + \Delta t) = \mathbf{r}_i(t) + \Delta t \mathbf{v}_i$$
$$\mathbf{v}_i(t + \Delta t) = \mathbf{v}_i(t) + \Delta t \mathbf{a}_i$$

- ▶ The Velocity Verlet algorithm we used previously is only correct for forces which are not a function of the velocity, but the viscous force terms in SPH are.

# Outline

Introduction

About Hydrodynamics

General Overview

Balance Equations

Solution Methods

Grids

Smoothing Function

Resulting Force Laws

Selected Smoothing Function

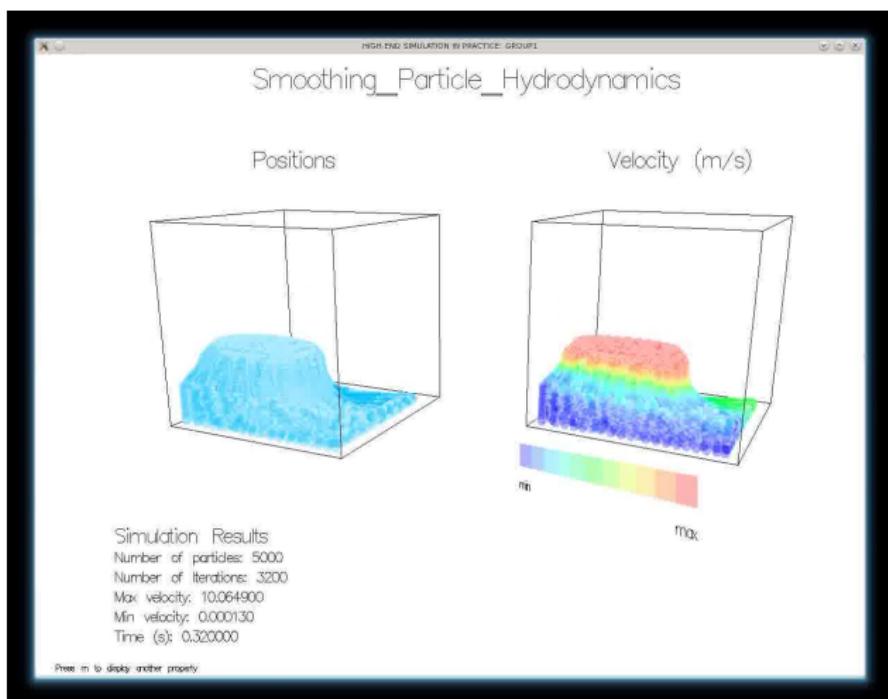
Pressure Force

Integrator

Showroom

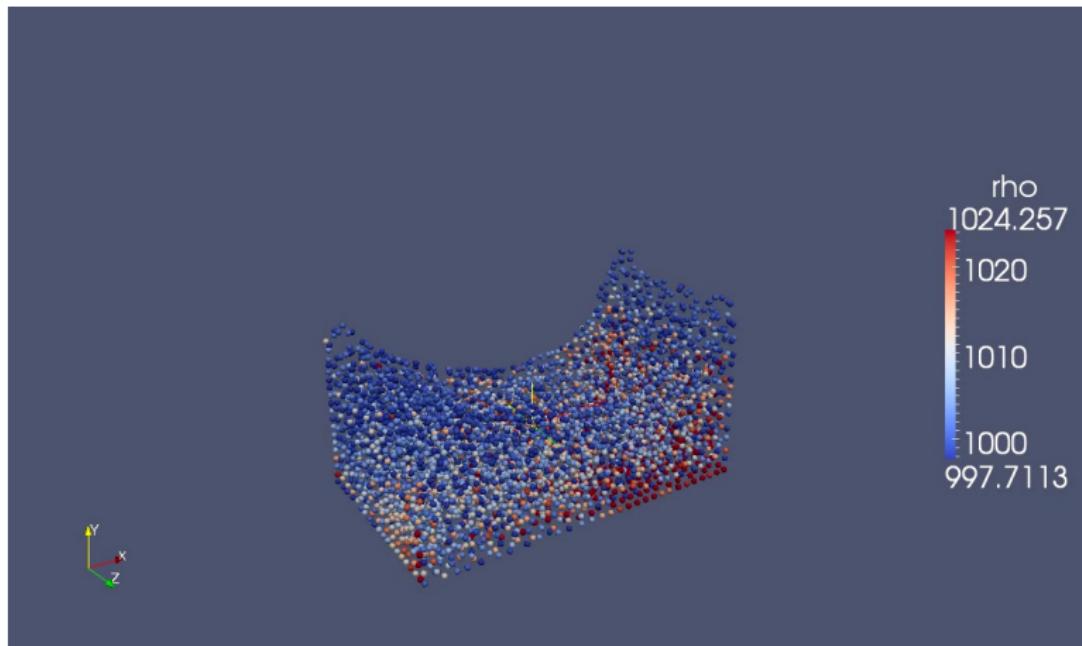
Final Remarks

# HESP Project SPH #1



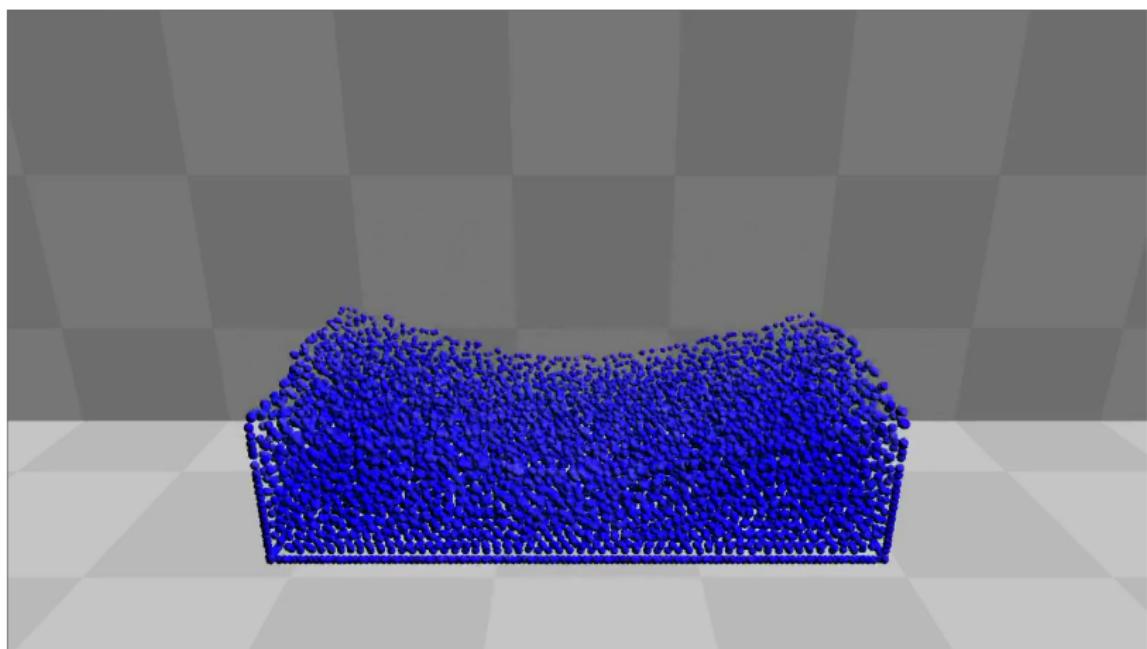
HESP 2013 - Group 1

# HESP Project SPH #1



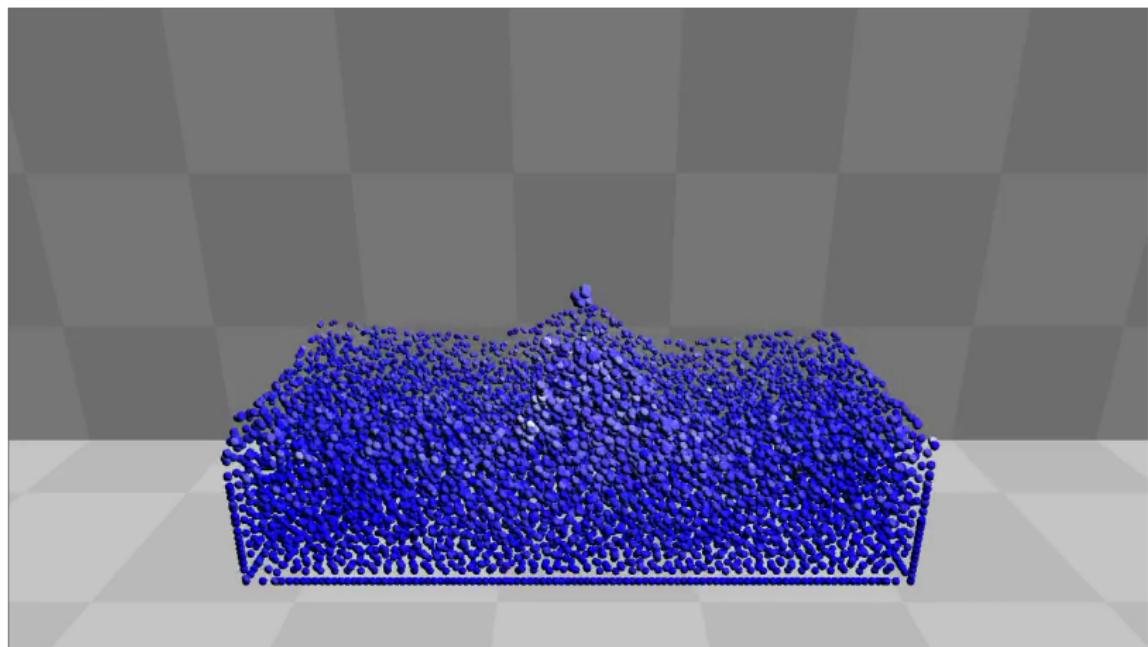
HESP 2013 - Group 1

## HESP Project SPH #2



HESP 2013 - Group 2

## HESP Project SPH #2



HESP 2013 - Group 2

# Outline

Introduction

About Hydrodynamics

General Overview

Balance Equations

Solution Methods

Grids

Smoothing Function

Resulting Force Laws

Selected Smoothing Function

Pressure Force

Integrator

Showroom

Final Remarks

# Summary

- ▶ Today you learned about neighbour lists, an important optimization for any MD implementation.
- ▶ In the next exercise you will extend your own MD implementation to make use of this.
- ▶ We showed you SPH, basically a modified MD with different force models.
- ▶ In the project phase one or more groups might actually implement SPH.